

# Lab 4: Out-of-Order PARCv2 Processor

Bryan SebaRaj

Professor Abhishek Bhattacharjee

EENG 420 - Computer Architecture

May 7, 2025

## Abstract

Out-of-order execution is a technique used in computer architecture design to improve the performance a processor by allowing instructions to be executed in a different order than they appear in the program. This lab focused on implementing an out-of-order PARCv2 processor with a reorder buffer and speculation. Splitting the implementation into two parts, I first implemented a reorder buffer to allow for out-of-order execution. I then extended the reorder buffer to support speculation, allowing the processor to execute instructions before it is known whether they will be needed. The processor was tested using a set of benchmarks, including vector add, complex multiply, masked filter, and binary search. The cycle counts and IPC for each benchmark were recorded and compared across the two implementations and a baseline pipelined processor with bypassing and no reorder buffer. The results showed that the out-of-order PARCv2 processors with a reorder buffer and speculation outperformed the baseline pipelined processor with bypassing in most cases, demonstrating the benefits of out-of-order execution and speculation.

## Design

### Objective 1: Out-of-Order PARCv2 Processor with Reorder Buffer

Given the existing I2O2 PARCv2 processor without a reorder buffer, the I2O1 PARCv2 processor was constructed by adding a reorder buffer, enabling bypassing out of the reorder buffer in the scoreboard, and modifying the datapath to enable values to be written from the reorder buffer instead of from the writeback stage. Specifically, the reorder buffer was built as a 16-entry circular queue, with a four-bit head pointer (commit pointer) and a four-bit tail pointer (allocate pointer). Each entry comprised a valid bit, a pending bit, and a five-bit field identifying the destination physical register.

Allocation is enabled whenever the buffer is not full, which is detected by `(head_ptr == tail_ptr) && valid[head_ptr]`, and, upon handshake, records the incoming physical register, asserting both valid and pending bits, and advancing the tail pointer. The fill port handles

writeback requests by deasserting the pending bit of the specified slot, marking that instruction's result as ready. Commit logic evaluates (`valid[head_ptr] && !pending[head_ptr]`) each cycle, and if true, deallocates the head entry, issues the register-file write-enable and corresponding write-address outputs, and increments the head pointer.

By separating the allocate, fill, and commit pathways and using simple pointer arithmetic, this should guarantee precise in-order retirement, support scoreboard-based bypassing of pending values, and remain deadlock-free and readily synthesizable.

## Objective 2: Out-of-Order PARCv2 Processor with Speculation

The speculative pv2-spec processor retains the original five-stage PARCv2 datapath but augments its ROB and control logic to allow instructions to issue and execute past unresolved branches. In my implementation, each of the 16 ROB entries has an additional one-bit "speculative" flag that is set at allocation whenever a branch is present in the Issue stage. All other fields, valid, pending, and the destination physical register, remain unchanged. In addition, I added four new control inputs to the ROB: a speculative-allocation indicator, and three branch-resolution signals (branch-resolved valid, resolved direction, and mispredict). When a branch retires in the Execute stage, the ROB scans all entries. On a correct prediction it simply clears each speculative bit and on a misprediction it invalidates every speculative entry.

Note that the commit logic was tightened so that by default only non-speculative, ready entries retire in-order, preventing speculative results from escaping until proven correct. By implementing a global invalidate or global clear approach, rather than precisely tracking the age of each branch, I minimize additional pointer arithmetic and per-entry state. In the CoreCtrl module, the original stall-on-branch mechanism was removed so that speculative instructions flow unimpeded. To address this, branch resolution now drives the ROB's speculative bits. Bypass support in the scoreboard and register file was reused unmodified from Part 1, since speculative and non-speculative values share the same forwarding paths.

Note that this implementation does not deviate from the prescribed datapath. No extensions were implemented.

## Testing Methodology

In order to test correctness, the pipelined processors were tested using the provided assembly test programs. These were designed to cover all the instructions in MIPS v2 assembly, including ALU operations, load and store operations, branch operations, jump operations, memory operations. In addition to correctness, performance was also measured with instructions per cycle and cycle counts through the provided benchmark test suite.

In testing the processor's functionality, I implemented four additional assembly tests that verified critical aspects of out-of-order and speculative execution. The commit-order-test created a specific dependency chain with a high-latency division operation followed by dependent instructions, validating that the ROB correctly maintained program order even when execution

completes out of sequence. The rob-bypass-test verified the processor's ability to bypass values from the ROB that have been written back but not yet committed, ensuring instruction-level parallelism doesn't compromise data integrity. The waw-test explores write after write hazards that execute correctly across processor implementations due to sufficient time separation, while the better-ipc-test revealed scenarios where the simpler bypassing processor achieved better performance than the complex out-of-order and speculative implementations due to ROB management overhead. This comprehensive test suite demonstrated the processors' speculative execution, data forwarding, and commitment mechanisms while targeting specific architectural vulnerabilities, providing confidence in the implementations' correctness across various dependency patterns. The assembly tests can be found in the `/tests/lab4-custom` directory.

## Evaluation

The pipelined in-order processor with bypassing, pipelined out-of-order processor with the re-order buffer, and pipelined out-of-order processor with speculation were tested using the following benchmarks: vector add, complex multiply, masked filter, and binary search. The cycle counts and IPC for each benchmark are shown below in the format, cycle count / IPC.

Benchmark	pv2byp	pv2ooo	pv2spec
binary search	1749/0.731275	1695/0.754572	2464/0.519075
complex multiply	15,312/0.121735	2562/0.727557	3138/0.594009
masked filter	13,832/0.32526	6248/0.720070	7693/0.584817
vector add	473/0.961945	513/0.886940	524/0.868321

## Discussion

Trivial analysis of the performance of the three processors with the provided benchmarks elucidates the significant benefits of out-of-order execution when compared to naive, in-order bypassing processor. The simple bypassed in-order pipeline (pv2byp) delivers excellent performance on straight-line code with minimal hazards (vector add IPC  $\approx 0.96$ ) but suffers on high-latency or data-parallel loops (complex multiply IPC  $\approx 0.12$ , masked filter IPC  $\approx 0.33$ ). Introducing a 16-entry ROB (pv2ooo) substantially increases instruction-level parallelism—raising IPC to 0.73 and 0.72 on those loops. But this is at the expense of slightly higher cycle time and modest pipeline control overhead, which yields little or even negative gain on trivial code (vector add IPC drops to 0.89). Further adding speculative execution (pv2spec) incurs additional misprediction and commit-logic overhead. It delivers only marginal benefit over pure out-of-order on predictable loops and dramatically degrades throughput on branch-heavy benchmarks (binary search IPC falls from 0.75 to 0.52). Thus, under the Iron Law ( $\text{CPU time} = \text{IC} \times \text{CPI} \times \text{cycle time}$ ), the ROB primarily lowers CPI when ILP is available, speculation can reduce CPI further only if branch prediction is accurate, and both techniques slightly increase cycle time without affecting instruction count. As such, their net benefit depends on workload characteristics.

## Figures

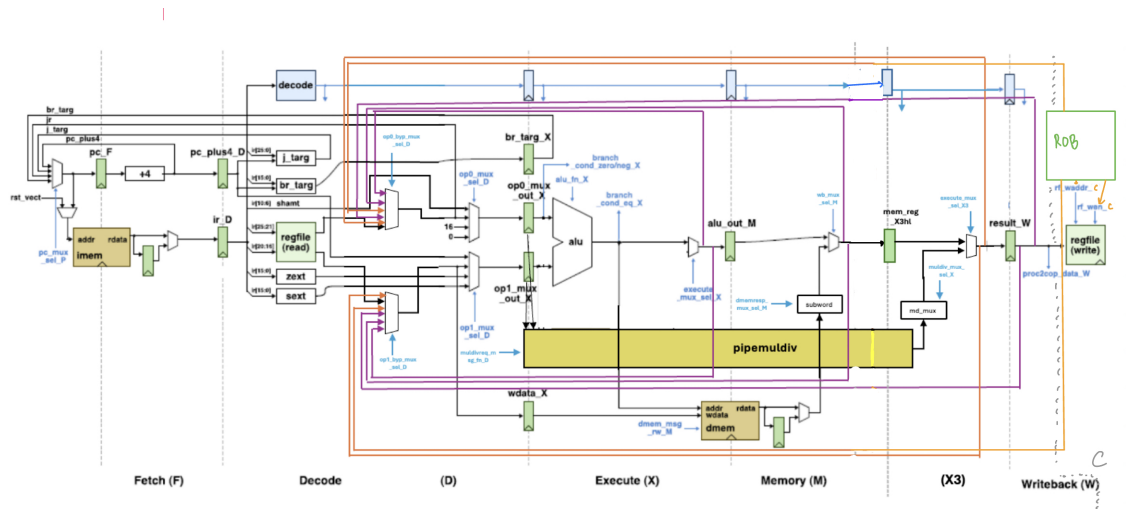


Figure 1. Datapath of out-of-order, single-issue processor.

Valid [6]	Pending [5]	Physical Register [4:0]
1	1	4'0011 ( $3_{10}$ )

Figure 2. Reorder buffer for part 1. Note that the data in the ROB (from the diagram provided on Canvas) live in the pipeline's writeback stage or go straight to the register file.

Speculative[7]	Valid [6]	Pending [5]	Physical Register [4:0]
0	1	1	4'0011 ( $3_{10}$ )

Figure 3. Reorder buffer for part 2. Note that the data in the ROB (from the diagram provided on Canvas) live in the pipeline's writeback stage or go straight to the register file.




Figure 4. Reorder buffer of out-of-order, single-issue processor.